

For this assignment, I found a repository called Star Ruler 2, and in this repository is a `str_util.cpp` file. This file contains methods to manipulate strings, e.g., `replace`, `split`, and `join`. I decided to test the methods in this file because they are pretty straight-forward and easy to fuzz. To test these methods, I created one big GTest called “StrManipulation”. I tested each individual method first to make sure the output was the same as `std::string` output. I then created a `OneOf`, so that I could test a random number of methods calls on a string. I also created helper methods that compare the `toUpper`, `toLowerCase`, and `replace` methods with equivalent `c` methods. I used the `DeepState` method `DeepState_CStr_C` to create a string from these characters:

“`abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ0123456789-_/?*%$#@![]{}`”. I wanted to see if any special characters would cause any unwanted behavior.

This package contains `split/join` methods to split a string into a vector, and then join that vector into a string. I would assume that you would be able to do a roundtrip and get the original string, but this does not seem the case. The test fails every time on any input. My guess is that I may be splitting on a char that is not in the string, but this should not cause any problems with the join. The method may not handle splitting on a character not in the string. I then commented the code out so that I could test the other methods.

I first used the `deepstate` `angr` fuzzer without a seed. It ran 24757 tests with 2 failing tests. My first run was with `-fuzz`, but this caused my output directory to be empty. The only way I could get failure to save was to save all output. The second time I ran it, I got 20480 passed tests and 1 failure. The end of the failure test’s output is:

```
TRACE: Tests.cpp(96): Number of Iterations: 9
TRACE: Tests.cpp(115): Started Trim: gtcg73eRReg]erkejsqf)ztG
TRACE: Tests.cpp(118): Trimmed whitespace
TRACE: Tests.cpp(96): Number of Iterations: 10
TRACE: Tests.cpp(128): Started Replace: gtcg73eRReg]erkejsqf)ztG
TRACE: Tests.cpp(143): Replaced substring: ]erkejsqf) with 4]2e@8TMYw string:gtcg73eRReg4]2e@8TMYwztG
TRACE: Tests.cpp(96): Number of Iterations: 11
TRACE: Tests.cpp(128): Started Replace: gtcg73eRReg4]2e@8TMYwztG
TRACE: Tests.cpp(143): Replaced substring: ]2e@8TMYwzt with r4x?XYMJT[K string:gtcg73eRReg4r4x?XYMJT[KG
TRACE: Tests.cpp(96): Number of Iterations: 12
TRACE: Tests.cpp(121): Started Char Trim: gtcg73eRReg4r4x?XYMJT[KG
TRACE: Tests.cpp(125): Trimmed char: string: gtcg73eRReg4r4x?XYMJT[KG
TRACE: Tests.cpp(96): Number of Iterations: 13
TRACE: Tests.cpp(121): Started Char Trim: gtcg73eRReg4r4x?XYMJT[KG
TRACE: Tests.cpp(125): Trimmed char: string: gtcg73eRReg4r4x?XYMJT[KG
TRACE: Tests.cpp(96): Number of Iterations: 14
TRACE: Tests.cpp(99): Started Lowercase: gtcg73eRReg4r4x?XYMJT[KG
TRACE: Tests.cpp(104): Our Lowercase: gtcg73erreg4r4x?xymjt[kg
TRACE: Tests.cpp(96): Number of Iterations: 15
TRACE: Tests.cpp(115): Started Trim: gtcg73erreg4r4x?xymjt[kg
TRACE: Tests.cpp(118): Trimmed whitespace
TRACE: Tests.cpp(96): Number of Iterations: 16
TRACE: Tests.cpp(99): Started Lowercase: gtcg73erreg4r4x?xymjt[kg
TRACE: Tests.cpp(104): Our Lowercase: gtcg73erreg4r4x?xymjt[kg
TRACE: Tests.cpp(96): Number of Iterations: 17
TRACE: Tests.cpp(115): Started Trim: gtcg73erreg4r4x?xymjt[kg
TRACE: Tests.cpp(118): Trimmed whitespace
TRACE: Tests.cpp(96): Number of Iterations: 18
TRACE: Tests.cpp(115): Started Trim: gtcg73erreg4r4x?xymjt[kg
TRACE: Tests.cpp(118): Trimmed whitespace
TRACE: Tests.cpp(96): Number of Iterations: 19
TRACE: Tests.cpp(121): Started Char Trim: gtcg73erreg4r4x?xymjt[kg
TRACE: Tests.cpp(125): Trimmed char: g string: tcg73erreg4r4x?xymjt[k
TRACE: Tests.cpp(96): Number of Iterations: 20
TRACE: Tests.cpp(128): Started Replace: tcg73erreg4r4x?xymjt[k
terminate called after throwing an instance of 'std::out_of_range'
what(): basic_string::substr: __pos (which is 23) > this->size() (which is 22)
ERROR: Failed: StrUtil_StrManipulation
ERROR: Test case angr_out/1e0bf949343f624237a2309124f034052101a4ae.fail failed
```

The output shows that the test failed is in the oneOf when trying to get a substring of a string. To get a substring, I am using `std::string::substr`. This is not a problem with my methods, but something that happened when picking indices for the substring to replace. Not 100% sure why this error occurred and why it occurs a very small amount of the time.

I decided to rerun the tests to see if something different happens, this time with `-exit_on_fail`. A crash occurred and the end of the output is:

```
TRACE: Tests.cpp(96): Number of Iterations: 12
TRACE: Tests.cpp(146): Started Replace Char: }Dm~-T4l*45 1mpU0j]E$N
TRACE: Tests.cpp(151): Replaced char: 42 with: 92 string: }Dm~-T4l\45 1mpU0j]E$N
TRACE: Tests.cpp(96): Number of Iterations: 13
TRACE: Tests.cpp(128): Started Replace: }Dm~-T4l\45 1mpU0j]E$N
TRACE: Tests.cpp(143): Replaced substring: ~-T with 7Tc string:}Dm7Tc4l\45 1mpU0j]E$N
TRACE: Tests.cpp(96): Number of Iterations: 14
TRACE: Tests.cpp(146): Started Replace Char: }Dm7Tc4l\45 1mpU0j]E$N
TRACE: Tests.cpp(151): Replaced char: 47 with: 61 string: }Dm7Tc4l\45 1mpU0j]E$N
TRACE: Tests.cpp(96): Number of Iterations: 15
TRACE: Tests.cpp(146): Started Replace Char: }Dm7Tc4l\45 1mpU0j]E$N
TRACE: Tests.cpp(151): Replaced char: -69 with: -16 string: }Dm7Tc4l\45 1mpU0j]E$N
TRACE: Tests.cpp(96): Number of Iterations: 16
TRACE: Tests.cpp(115): Started Trim: }Dm7Tc4l\45 1mpU0j]E$N
TRACE: Tests.cpp(118): Trimmed whitespace
TRACE: Tests.cpp(96): Number of Iterations: 17
TRACE: Tests.cpp(107): Started Uppercase: }Dm7Tc4l\45 1mpU0j]E$N
TRACE: Tests.cpp(112): Our Uppercase: }DM7TC4L\45 1MPU0J]E$N
TRACE: Tests.cpp(96): Number of Iterations: 18
TRACE: Tests.cpp(128): Started Replace: }DM7TC4L\45 1MPU0J]E$N
```

This test crashed in the same spot, but this time a timeout occurred, not a failure. The timeout occurred in the `replace()` function we are testing. Again, I am not sure why this test timed out.

To end my angr testing, I decided to try pure swarm testing. This resulted in 5851 passing tests and 1 failure. The failure occurred for the same reason as the failure above. The output would not save though.

I then ran DeepState with AFL and used the same string as above as the input. The output is:

american fuzzy lop 2.52b (string_AFL)			
process timing		overall results	
run time : 0 days, 1 hrs, 49 min, 59 sec		cycles done : 1	
last new path : 0 days, 1 hrs, 6 min, 57 sec		total paths : 343	
last uniq crash : 0 days, 1 hrs, 11 min, 13 sec		uniq crashes : 39	
last uniq hang : 0 days, 1 hrs, 38 min, 37 sec		uniq hangs : 1	
cycle progress		map coverage	
now processing : 168* (48.98%)		map density : 1.35% / 1.80%	
paths timed out : 0 (0.00%)		count coverage : 3.34 bits/tuple	
stage progress		findings in depth	
now trying : auto extras (over)		favored paths : 31 (9.04%)	
stage execs : 2590/9800 (26.43%)		new edges on : 57 (16.62%)	
total execs : 1.42M		total crashes : 441 (39 unique)	
exec speed : 719.1/sec		total tmouts : 1801 (43 unique)	
fuzzing strategy yields		path geometry	
bit flips : 21/40.3k, 2/40.2k, 7/40.1k		levels : 5	
byte flips : 0/5033, 0/4625, 2/4563		pending : 292	
arithmetics : 2/259k, 0/85.5k, 0/41.9k		pend fav : 0	
known ints : 0/24.4k, 4/108k, 2/181k		own finds : 342	
dictionary : 0/0, 0/0, 1/161k		imported : n/a	
havoc : 340/419k, 0/0		stability : 100.00%	
trim : 7.99%/1962, 7.23%			
[cpu000:214%]			

The AFL test ran for almost 2 hours (without -d), completing one cycle and finding 343 paths. I was not surprised by the path count since there is a lot of randomization in the test, i.e., replacing a random substring each time. However, I was surprised by the number of failures and timeouts. There was a total of 441 crashes, 39 unique, and 1801 timeouts, 43 unique. I reran some of the tests, and they are either timeouts or failures because of the `std::out_of_range` exception.

I then tried running AFL with pure swarm testing. This resulted in the following:

american fuzzy lop 2.52b (string_AFL_swarm)			
process timing		overall results	
run time : 0 days, 8 hrs, 2 min, 42 sec		cycles done : 5	
last new path : 0 days, 1 hrs, 45 min, 20 sec		total paths : 451	
last uniq crash : 0 days, 0 hrs, 13 min, 19 sec		uniq crashes : 42	
last uniq hang : 0 days, 6 hrs, 3 min, 6 sec		uniq hangs : 1	
cycle progress		map coverage	
now processing : 424* (94.01%)		map density : 1.21% / 1.92%	
paths timed out : 0 (0.00%)		count coverage : 3.97 bits/tuple	
stage progress		findings in depth	
now trying : bitflip 2/1		favored paths : 40 (8.87%)	
stage execs : 10.0k/65.6k (15.26%)		new edges on : 58 (12.86%)	
total execs : 13.1M		total crashes : 7659 (42 unique)	
exec speed : 592.7/sec		total tmouts : 4983 (56 unique)	
fuzzing strategy yields		path geometry	
bit flips : 25/821k, 9/755k, 8/755k		levels : 9	
byte flips : 1/94.5k, 2/46.7k, 1/47.2k		pending : 145	
arithmetics : 3/2.59M, 1/932k, 0/337k		pend fav : 0	
known ints : 0/244k, 2/1.14M, 14/1.96M		own finds : 450	
dictionary : 0/0, 0/0, 8/2.27M		imported : n/a	
havoc : 418/1.04M, 0/0		stability : 100.00%	
trim : 23.70%/32.6k, 50.39%			
[cpu000:193%]			

I let AFL run for 8 hours without the -d flag. It went through 5 cycles and found 451 paths. To my surprise, 42 unique crashes and 56 unique timeouts occurred, but the total number of crashes and timeouts is very large. As I watched it at the beginning, I noticed that it also took a little longer for AFL to find the crashes. This could have been because the replace method was being called less (not being used a lot at the beginning).

I tried using LibFuzzer but could not get it to work. If I used deepstate-libfuzzer, I got a python arg parse error. I tested this with the RunLen DeepState example. I also tried running the executable directly, but it also failed to run with "Killing the process and childs."

I also could not get Eclipse to run either. When I tried to run my executable with deepstate-eclisper, I get an error about my output directory not being there, when it clearly is. I tried to give a full path, but this did not fix the problem.

In order to use Honggfuzz, I had to use the new DeepState Docker container. I let the fuzzer run for almost 9 hours. The end of the output is below:

```
FUZZ_STATS:deepstate.core.fuzz:-----
FUZZ_STATS:deepstate.core.fuzz|unique_crashes:1
FUZZ_STATS:deepstate.core.fuzz:fuzzer_pid:1264
FUZZ_STATS:deepstate.core.fuzz:start_time:1583820948
FUZZ_STATS:deepstate.core.fuzz:-----
FUZZ_STATS:deepstate.core.fuzz|unique_crashes:1
FUZZ_STATS:deepstate.core.fuzz:fuzzer_pid:1264
FUZZ_STATS:deepstate.core.fuzz:start_time:1583820948
FUZZ_STATS:deepstate.core.fuzz:-----
FUZZ_STATS:deepstate.core.fuzz|unique_crashes:1
FUZZ_STATS:deepstate.core.fuzz:fuzzer_pid:1264
FUZZ_STATS:deepstate.core.fuzz:start_time:1583820948
FUZZ_STATS:deepstate.core.fuzz:-----
FUZZ_STATS:deepstate.core.fuzz|unique_crashes:1
FUZZ_STATS:deepstate.core.fuzz:fuzzer_pid:1264
FUZZ_STATS:deepstate.core.fuzz:start_time:1583820948
FUZZ_STATS:deepstate.core.fuzz:-----
FUZZ_STATS:deepstate.core.fuzz|unique_crashes:1
FUZZ_STATS:deepstate.core.fuzz:fuzzer_pid:1264
FUZZ_STATS:deepstate.core.fuzz:start_time:1583820948
FUZZ_STATS:deepstate.core.fuzz:-----
FUZZ_STATS:deepstate.core.fuzz|unique_crashes:1
FUZZ_STATS:deepstate.core.fuzz:fuzzer_pid:1264
FUZZ_STATS:deepstate.core.fuzz:start_time:1583820948
FUZZ_STATS:deepstate.core.fuzz:-----
FUZZ_STATS:deepstate.core.fuzz|unique_crashes:1
FUZZ_STATS:deepstate.core.fuzz:fuzzer_pid:1264
FUZZ_STATS:deepstate.core.fuzz:start_time:1583820948
FUZZ_STATS:deepstate.core.fuzz:-----
FUZZ_STATS:deepstate.core.fuzz|unique_crashes:1
FUZZ_STATS:deepstate.core.fuzz:fuzzer_pid:1264
FUZZ_STATS:deepstate.core.fuzz:start_time:1583820948
FUZZ_STATS:deepstate.core.fuzz:-----
```

starRuler (bash)

XCI user@Od24af8e40ad: ~/kdi222\_cs499\_sp20/assignment2/starRuler... XCI

user@c6c524bf6926: ~/deepstate/kdi222\_cs499\_sp20/assignment2 XCI

Honggfuzz found 1 unique crash. After replaying the crash, I found that the crash was caused by the same `std::string::out_of_range` error.

Overall, I am not sure if the methods I tested caused any of the failures, but if how I wrote the tests caused them. Most of the time, they pass, but sometimes a `std::out_of_range` exception occurs or a timeout occurs in the `replace()` method. However, the timeout may actually be a real problem.